

# Tail Calling Between Code Generated by C and Native Backends

Laurent Huberdeau / Marc Feeley



# Gambit Compiler Pipeline

Scheme

GVM

C backend

JS backend

C code

JavaScript code

```
(define (main)
  (println (f 9)))

(define (f x)
  (square (fx+ x 1)))

(define (square y)
  (fx* y y))
```

frontend

```
**** #<procedure main> =
#1 fs=0 entry-point nparams=0 ()
  frame[1] = r0
  r1 = '9
  jump fs=4 global[f] r0=#2 nargs=1
#2 fs=4 return-point
  r0 = frame[1]
  jump fs=0 global[println] nargs=1

**** #<procedure f> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx+ r1 '1)
  jump fs=0 global[square] nargs=1

**** #<procedure square> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx* r1 r1)
  jump fs=0 r0
```

# Gambit Compiler Pipeline

Scheme

GVM

```
(define (main)
  (println (f 9)))

(define (f x)
  (square (fx+ x 1)))

(define (square y)
  (fx* y y))
```

frontend

```
**** #<procedure main> =
#1 fs=0 entry-point nparams=0 ()
  frame[1] = r0
  r1 = '9
  jump fs=4 global[f] r0=#2 nargs=1
#2 fs=4 return-point
  r0 = frame[1]
  jump fs=0 global[println] nargs=1

**** #<procedure f> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx+ r1 '1)
  jump fs=0 global[square] nargs=1

**** #<procedure square> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx* r1 r1)
  jump fs=0 r0
```

C backend

C code

JS backend

JavaScript code

Native backend  
(x86-32/64, ARM)

Native code

NEW!

# C/Native Interoperability

## Scheme

## GVM

```
(define (main)
  (println (f 9)))

(define (f x)
  (square (fx+ x 1)))
```

frontend

```
**** #<procedure main> =
#1 fs=0 entry-point nparams=0 ()
  frame[1] = r0
  r1 = '9
  jump fs=4 global[f] r0=#2 nargs=1
#2 fs=4 return-point
  r0 = frame[1]
  jump fs=0 global[println] nargs=1

**** #<procedure f> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx+ r1 '1)
  jump fs=0 global[square] nargs=1
```

C backend

C code

```
(define (square y)
  (fx* y y))
```

frontend

```
**** #<procedure square> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx* r1 r1)
  jump fs=0 r0
```

Native backend  
(x86-32/64, ARM)

Native code

Separate compilation of modules by  
C and Native backends

# C/Native Interoperability

## Scheme

## GVM

```
(define (main)
  (println (f 9)))

(define (f x)
  (square (fx+ x 1)))
```

frontend

```
**** #<procedure main> =
#1 fs=0 entry-point nparams=0 ()
  frame[1] = r0
  r1 = '9
  jump fs=4 global[f] r0=#2 nargs=1
#2 fs=4 return-point
  r0 = frame[1]
  jump fs=0 global[println] nargs=1

**** #<procedure f> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx+ r1 '1)
  jump fs=0 global[square] nargs=1
```

C backend

C code

```
(define (square y)
  (fx* y y))
```

frontend

```
**** #<procedure square> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx* r1 r1)
  jump fs=0 r0
```

Native backend  
(x86-32/64, ARM)

Native code

CROSS  
JUMPS

# C/Native Interoperability

## Scheme

## GVM

```
(define (main)
  (println (f 9)))

(define (f x)
  (square (fx+ x 1)))
```

frontend

```
**** #<procedure main> =
#1 fs=0 entry-point nparams=0 ()
  frame[1] = r0
  r1 = '9
  jump fs=4 global[f] r0=#2 nargs=1
#2 fs=4 return-point
  r0 = frame[1]
  jump fs=0 global[println] nargs=1

**** #<procedure f> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx+ r1 '1)
  jump fs=0 global[square] nargs=1
```

C backend

C code

```
(define (square y)
  (fx* y y))
```

frontend

```
**** #<procedure square> =
#1 fs=0 entry-point nparams=1 ()
  r1 = (fx* r1 r1)
  jump fs=0 r0
```

Native backend  
(x86-32/64, ARM)

Native code

How?

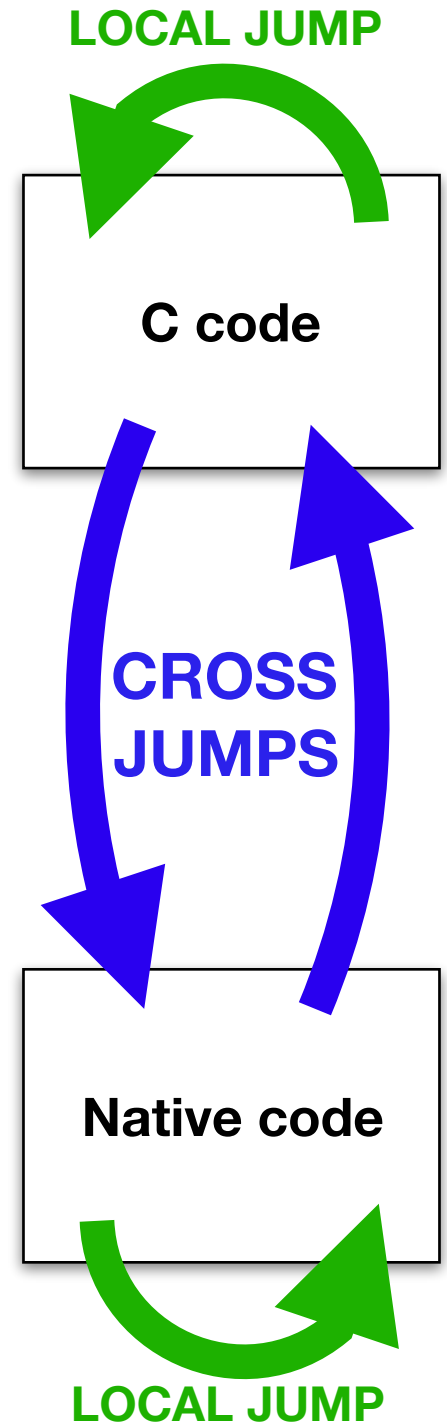
CROSS  
JUMPS

# Why?

- Reuse runtime library of C backend
  - Complex C functions (garbage collector, OS interface, ...)
  - Complex Scheme library procedures (`eval`, bignums, ...)
  - Complex I/O procedures (`pretty-print`, `read`, ...)
  - Simple primitives, like `##cons` and `##flsqrt` that are not yet inlined by Native backend (useful during development of backend)
- Mix modules compiled by C and Native backends and compile performance critical parts with Native backend

# Desiderata

1. **Cross jumps are transparent** (don't know ahead of time if a jump is going to be a cross jump or a local jump)
2. **Native code jumps are plain *branch-to-address* instructions** (no overhead on native local jumps due to support for cross jumps)
3. **Cost of cross jumps is similar to C backend local jumps** (which are performed with a trampoline)





# C Backend

- Uses a **processor state structure** to store the state of the Gambit Virtual Machine (registers, program counter, ...)

```
typedef struct processor_state_struct {  
    WORD r0, r1, r2, r3, r4;    /* registers    */  
    WORD *fp;                   /* frame ptr   */  
    WORD *hp;                   /* heap ptr    */  
    WORD *pc;                   /* program ctr */  
    ...  
} processor_state_struct;
```

- Uses a **label structure** to represent control points in the code (either procedure entry point or return point)


```
typedef struct label_struct {  
    WORD HEADER;  
    WORD entry;  
    WORD unused;  
    void (*host)(processor_state_struct *ps);  
} label_struct;
```

# C Backend

- Uses a **processor state structure** to store the state of the Gambit Virtual Machine (registers, program counter, ...)

```
typedef struct processor_state_struct {  
    WORD r0, r1, r2, r3, r4; /* registers */  
    WORD *fp; /* frame ptr */  
    WORD *hp; /* heap ptr */  
    WORD *pc; /* program ctr */  
    ...  
} processor_state_struct;
```

pc points to the  
currently executing  
control point



- Uses a **label structure** to represent control points in the code (either procedure entry point or return point)

```
typedef struct label_struct {  
    WORD HEADER;  
    WORD entry;  
    WORD unused;  
    void (*host)(processor_state_struct *ps);  
} label_struct;
```

# C Backend

- Uses a **processor state structure** to store the state of the Gambit Virtual Machine (registers, program counter, ...)

```
typedef struct processor_state_struct {  
    WORD r0, r1, r2, r3, r4; /* registers */  
    WORD *fp; /* frame ptr */  
    WORD *hp; /* heap ptr */  
    WORD *pc; /* program ctr */  
    ...  
} processor_state_struct;
```

pc points to the  
currently executing  
control point

- Uses a **label structure** to represent control points in the code (either procedure entry point or return point)

```
typedef struct label_struct {  
    WORD HEADER;  
    WORD entry;  
    WORD unused;  
    void (*host) (processor_state_struct *ps);  
} label_struct;
```

The host function is  
the C function  
containing this  
control point

# Example

```
(define (main) (println (f 9)))
```

```
(define (f x) (square (fx+ x 1)))
```

```
(define (square y) (fx* y y))
```

# Example

Non-tail call

```
(define (main) (println (f 9)))
```

```
(define (f x) (square (fx+ x 1)))
```

```
(define (square y) (fx* y y))
```

Tail call

```

void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2;                /* tagged fixnum = 9      */
                ps->r0 = 1+(WORD)&labels[1];    /* tagged return address */
                ps->pc = glo_f;                /* non-tail call f       */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);
    }
}

```

```

void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4;                    /* add 1 to arg          */
                ps->pc = glo_square;          /* tail call square     */
                return;
    }
}

```

```

void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2);        /* square arg            */
                ps->pc = ps->r0;                /* return to main       */
                return;
    }
}

```

C host functions

```

void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2; /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f; /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);

    }
}

```

**Basic Blocks**  
(first-class control points)

```

void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4; /* add 1 to arg */
                ps->pc = glo_square; /* tail call square */
                return;

    }
}

```

```

void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2); /* square arg */
                ps->pc = ps->r0; /* return to main */
                return;

    }
}

```

```

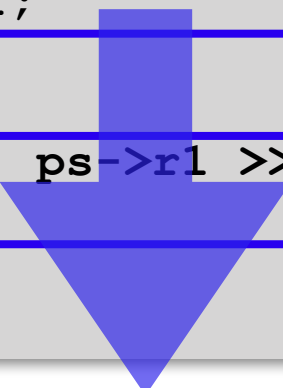
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2; /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f; /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);

    }
}

```



```

ps->r1 = 9 << 2;
ps->r0 = 1+(WORD)&labels[1];
ps->pc = glo_f;

```

```

(define (main) (println (f 9)))

```

```

(define (f x) (square (fx+ x 1)))

```



```

void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2; /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f; /* tagged return address */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);
    }
}

```

Return address for call to f refers to second basic block

```

ps->r1 = 9 << 2;
ps->r0 = 1+(WORD)&labels[1];
ps->pc = glo_f;

```

```

(define (main) (println (f 9)))

```

```

(define (f x) (square (fx+ x 1)))

```

# Control Points

```
label_struct labels[] = {  
    { ..., ..., ..., host_main } /* 0 */  
    , { ..., ..., ..., host_main } /* 1 */  
    , { ..., ..., ..., host_f } /* 2 */  
    , { ..., ..., ..., host_square } /* 3 */  
};
```

```
WORD glo_main      = 1+ (WORD) &labels[0];  
WORD glo_f         = 1+ (WORD) &labels[2];  
WORD glo_square    = 1+ (WORD) &labels[3];
```

# Control Points

```
label_struct labels[] = {  
    { ..., ..., ..., host_main } /* 0 */  
    , { ..., ..., ..., host_main } /* 1 */  
    , { ..., ..., ..., host_f } /* 2 */  
    , { ..., ..., ..., host_square } /* 3 */  
};
```

```
WORD glo_main      = 1+ (WORD) &labels[0];  
WORD glo_f         = 1+ (WORD) &labels[2];  
WORD glo_square    = 1+ (WORD) &labels[3];
```

Procedures and return addresses are tagged

Return addresses are first-class objects in Gambit Scheme  
(allows continuation inspection in Scheme and simplifies GC)

```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main },
    , { ..., ..., ..., host_main },
    , { ..., ..., ..., host_f },
    , { ..., ..., ..., host_square },
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2;                /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1];      /* tagged return address */
                ps->pc = glo_f;                  /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);

    }
}
```

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {


        case 0: ps->r1 += 4;                      /* add 1 to arg */
                ps->pc = glo_square;              /* tail call square */
                return;

    }
}
```

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2);          /* square arg */
                ps->pc = ps->r0;                  /* return to main */
                return;

    }
}
```



```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main }, ← ps->pc
    , { ..., ..., ..., host_main }
    , { ..., ..., ..., host_f }
    , { ..., ..., ..., host_square }
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

Execution starts with call to trampoline  
with ps->pc referring to program entry point

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2;                /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1];    /* tagged return address */
                ps->pc = glo_f;                /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);

    }
}
```

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4;                    /* add 1 to arg */
                ps->pc = glo_square;          /* tail call square */
                return;

    }
}
```

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2);        /* square arg */
                ps->pc = ps->r0;                /* return to main */
                return;

    }
}
```

1

```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main } ← ps->pc
    , { ..., ..., ..., host_main }
    , { ..., ..., ..., host_f }
    , { ..., ..., ..., host_square }
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

2

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {
        case 0: ps->r1 = 9 << 2; /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f; /* non-tail call f */
                return;
        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);
    }
}
```

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {
        case 0: ps->r1 += 4; /* add 1 to arg */
                ps->pc = glo_square; /* tail call square */
                return;
    }
}
```

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {
        case 0: ps->r1 *= (ps->r1 >> 2); /* square arg */
                ps->pc = ps->r0; /* return to main */
                return;
    }
}
```

1

```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main },
    , { ..., ..., ..., host_main } ← ps->r0
    , { ..., ..., ..., host_f } ← ps->pc
    , { ..., ..., ..., host_square }
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

2

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2;           /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f;             /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);
    }
}
```

3

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4;           /* add 1 to arg */
                ps->pc = glo_square;  /* tail call square */
                return;
    }
}
```

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2); /* square arg */
                ps->pc = ps->r0;         /* return to main */
                return;
    }
}
```

1

```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main },
    , { ..., ..., ..., host_main } ← ps->r0
    , { ..., ..., ..., host_f },
    , { ..., ..., ..., host_square } ← ps->pc
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

2

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2;           /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f;             /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);

    }
}
```

3

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4;           /* add 1 to arg */
                ps->pc = glo_square;  /* tail call square */
                return;

    }
}
```

4

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2); /* square arg */
                ps->pc = ps->r0;         /* return to main */
                return;

    }
}
```



1

```
void trampoline(processor_state_struct *ps) {
    while (TRUE) GET_HOST(ps->pc) (ps);
}
```

```
label_struct labels[] = {
    { ..., ..., ..., host_main },
    , { ..., ..., ..., host_main } ← ps->pc
    , { ..., ..., ..., host_f },
    , { ..., ..., ..., host_square }
};
```

```
WORD glo_main    = 1+(WORD)&labels[0];
WORD glo_f       = 1+(WORD)&labels[2];
WORD glo_square  = 1+(WORD)&labels[3];
```

2

```
void host_main(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[0])) / sizeof(label_struct)) {

        case 0: ps->r1 = 9 << 2; /* tagged fixnum = 9 */
                ps->r0 = 1+(WORD)&labels[1]; /* tagged return address */
                ps->pc = glo_f; /* non-tail call f */
                return;

        case 1: printf("%d\n", ps->r1 >> 2);
                exit(0);
    }
}
```

3

```
void host_f(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[2])) / sizeof(label_struct)) {

        case 0: ps->r1 += 4; /* add 1 to arg */
                ps->pc = glo_square; /* tail call square */
                return;
    }
}
```

4

```
void host_square(processor_state_struct *ps) {
    switch ((ps->pc - (1+(WORD)&labels[3])) / sizeof(label_struct)) {

        case 0: ps->r1 *= (ps->r1 >> 2); /* square arg */
                ps->pc = ps->r0; /* return to main */
                return;
    }
}
```

# Trampoline Optimizations

1. Locality is exploited by replacing **return** with a **goto** back to the **switch**
2. In *single host* compilation mode, all the code of a compiled file is generated in a single host function (exploits locality of jumps within the file)
3. Computed **goto**, available in gcc and clang, is faster than using a **switch**
4. Hosts can cache part of the processor state in local variables (improves likelihood C compiler will use machine registers)

# Native Backend

- Uses machine registers to store the state of the GVM :

```
rdi    ;; r0
rax    ;; r1
rbx    ;; r2
rdx    ;; r3
rsi    ;; r4
rsp    ;; frame pointer
rbp    ;; heap pointer
rcx    ;; processor state pointer
```

x86-64

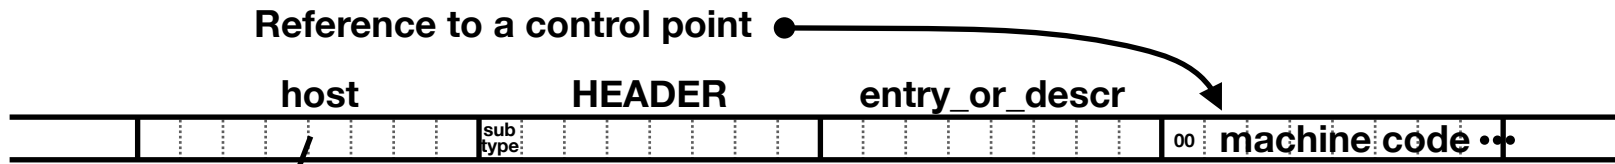
- This assignment of machine registers is different than the one used by the C compiler (which we can't and shouldn't control because it would interfere with the optimizations of the C compiler)

# Our Approach : Bridges and Dual Purpose `label_struct`

- Use handlers (**bridges**) for cross jumps that move the GVM state where the destination code expects it
- For Native backend : place executable machine code in `label_structs` starting at byte pointed by control point reference (to allow plain *branch-to-address* instructions in native code)
- Both backends use `label_structs` to represent control points :
  - Native backend : host field in `label_struct` points to the `to_native_bridge`
  - C backend : executable machine code in `label_struct` jumps to the `from_native_bridge`

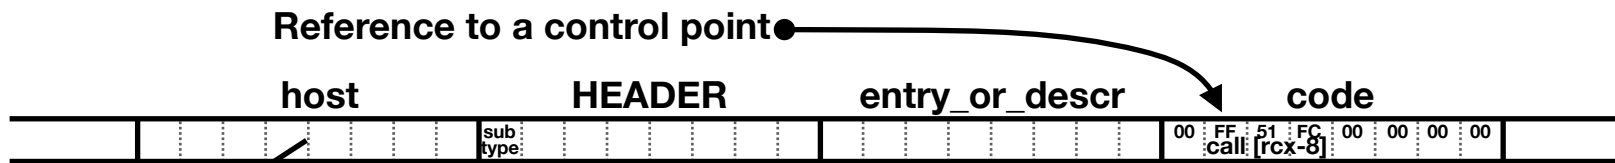
# label\_struct Layouts

x86-64 backend



```
void to_native_bridge(processor_state_struct *ps) { ... }
```

C backend

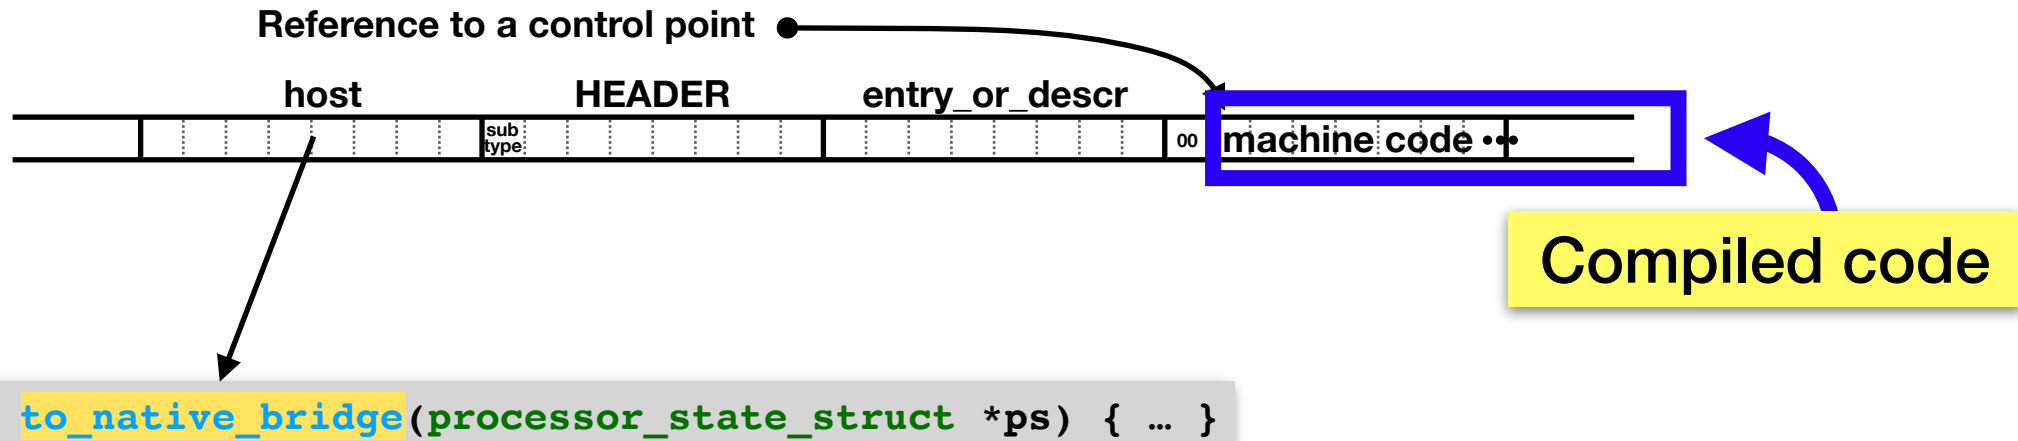


```
void host(processor_state_struct *ps) { ... }
```

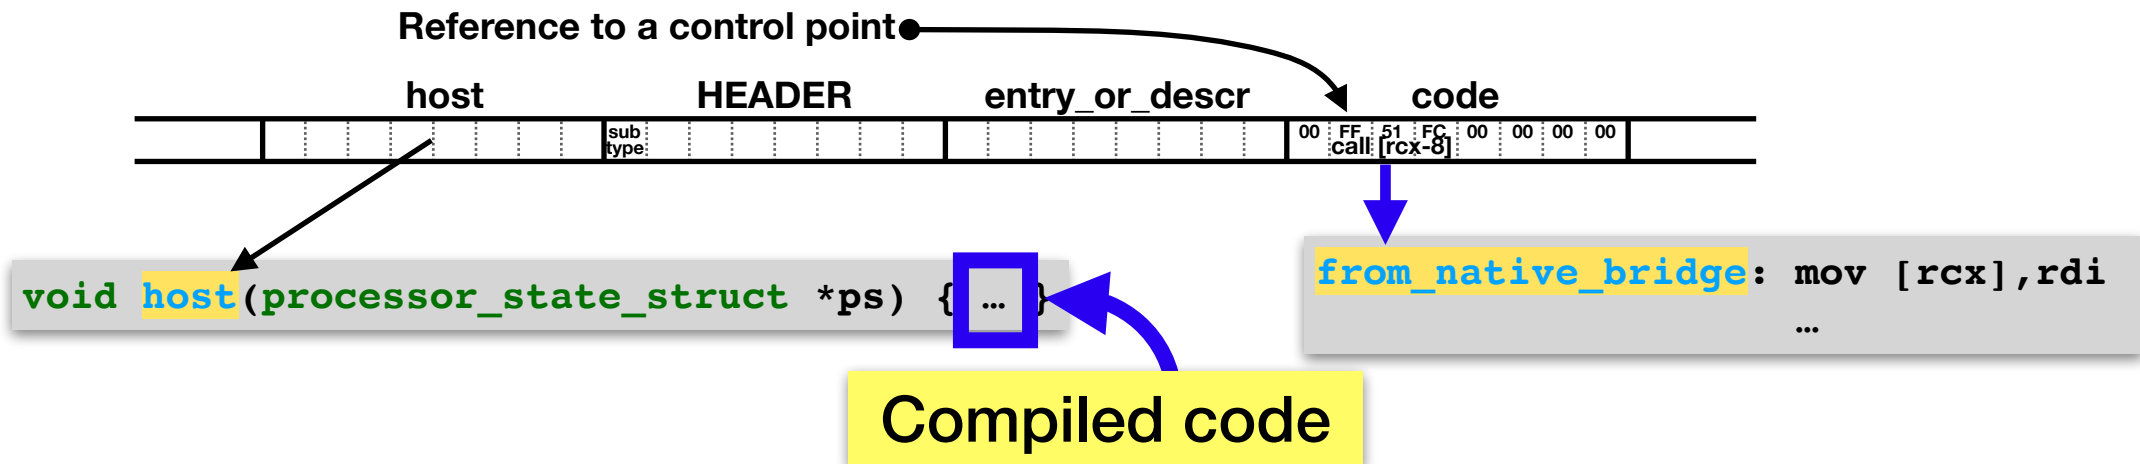
```
from_native_bridge: mov [rcx], rdi  
...
```

# label\_struct Layouts

x86-64 backend

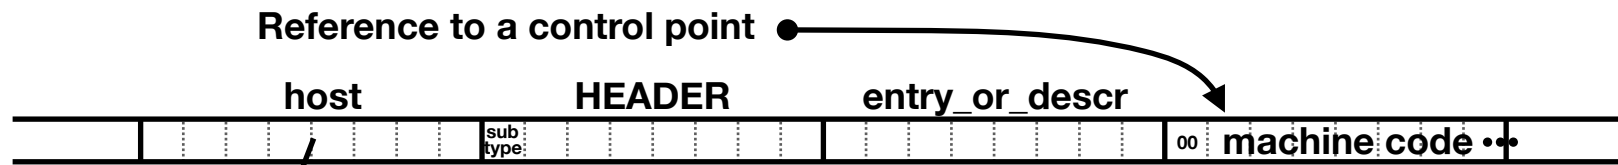


C backend



# label\_struct Layouts

x86-64 backend

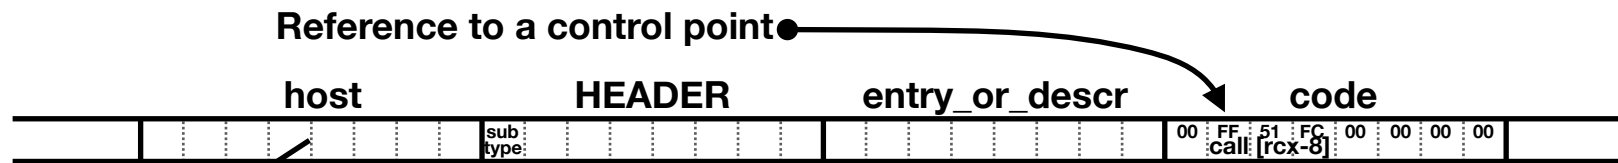


```
void to_native_bridge(processor_state_struct *ps) { ... }
```

Jumps to the machine code after storing the GVM state to registers

Bridge called on C to Native cross jump

C backend

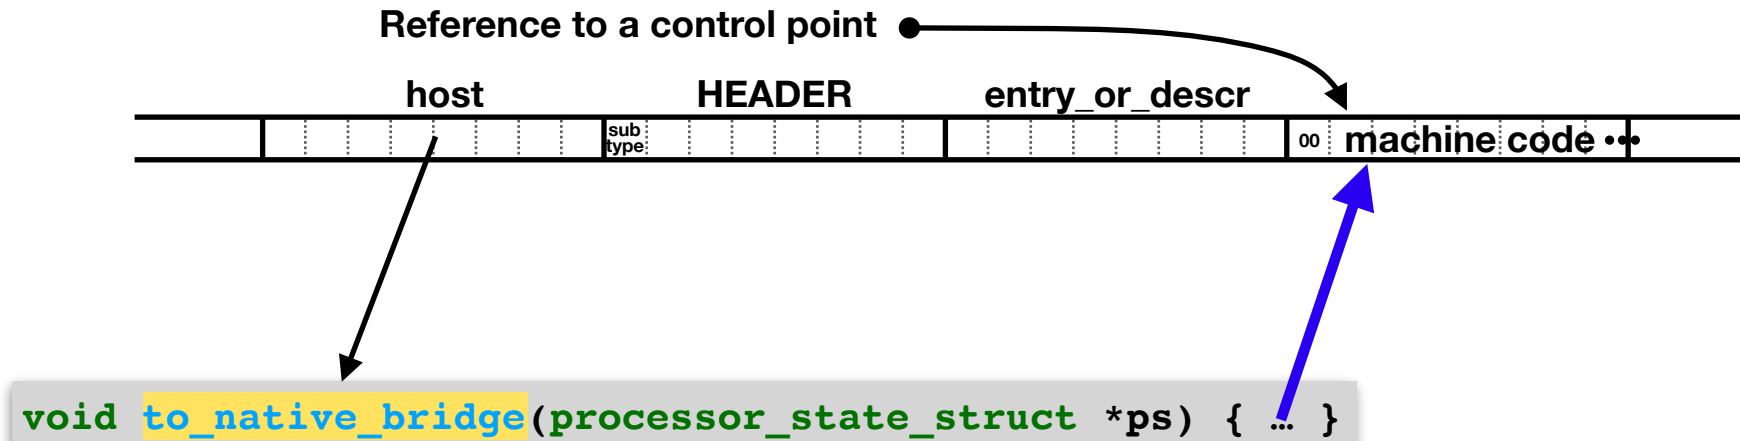


```
void host(processor_state_struct *ps) { ... }
```

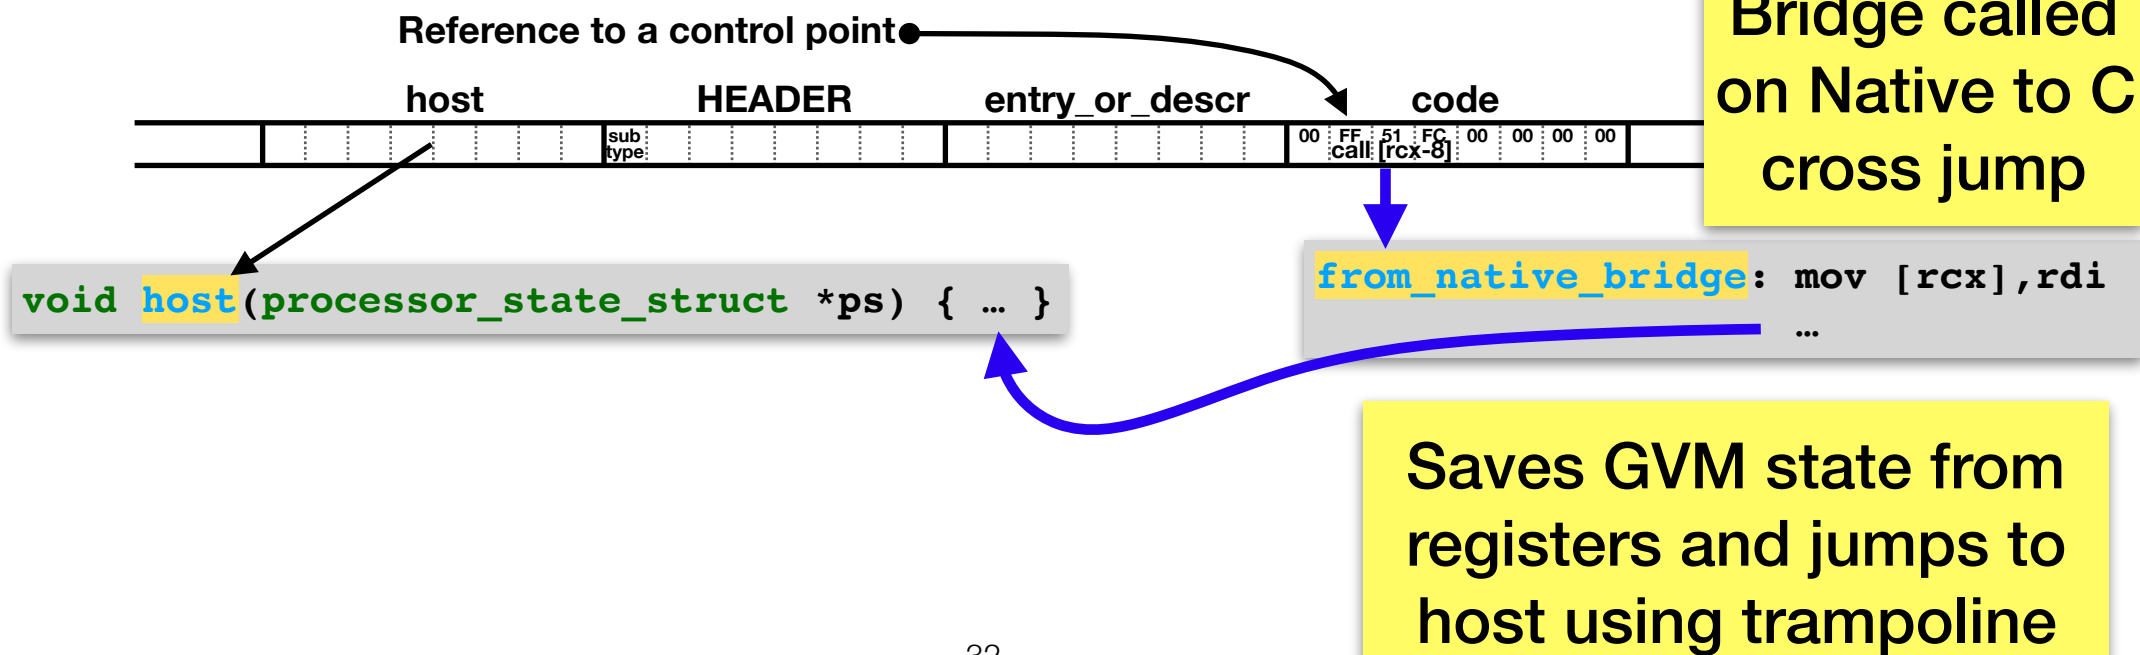
```
from_native_bridge: mov [rcx], rdi
...
```

# label\_struct Layouts

x86-64 backend



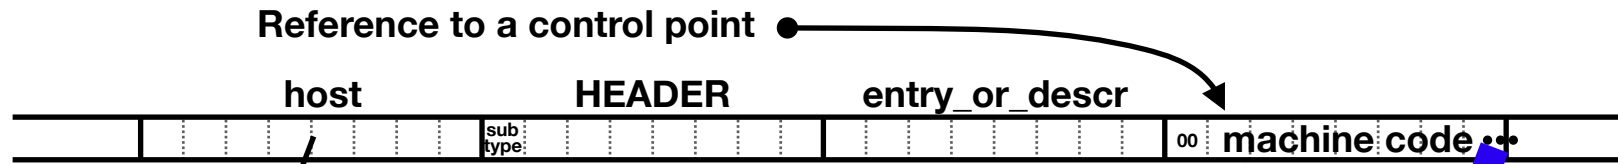
C backend





# label\_struct Layouts

x86-64 backend

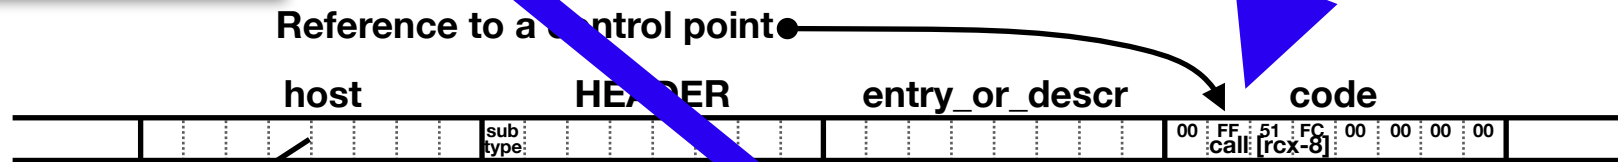


```
void to_native_bridge(processor_state_struct *ps) { ... }
```

Native to C  
cross jump

C to Native  
cross jump

C backend



```
void host(processor_state_struct *ps) { ... }
```

```
from_native_bridge: mov [rcx], rdi
...
```

# Implementation of Bridges

x86-64

```
void to_native_bridge(processor_state_struct *ps) {
    __asm__ __volatile__ (

        "mov  %0, %%rcx" // copy ps into rcx

        "mov  %%rsp, -2*8(%%rcx)" // save C sp

        // setup handler for returning from native code
        "lea  from_native_bridge(%%rip), %%rax"
        "mov  %%rax, -1*8(%%rcx)" // setup handler

        // setup frame pointer and heap pointer registers
        "mov  5*8(%%rcx), %%rsp" // rsp = ps->fp
        "mov  6*8(%%rcx), %%rbp" // rbp = ps->hp

        // setup self register
        "mov  4*8(%%rcx), %%rsi" // rsi = ps->r4

        "mov  8*8(%%rcx), %%rax" // rax = ps->pc
        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  setup_other_registers"
        "add  $3, %%rsi" // handle closures
        "push %%rsi"
        "add  $-3, %%rsi"

        "setup_other_registers:"
        "mov  (%%rcx), %%rdi" // rdi = ps->r0
        "mov  1*8(%%rcx), %%rax" // rax = ps->r1
        "mov  2*8(%%rcx), %%rbx" // rbx = ps->r2
        "mov  3*8(%%rcx), %%rdx" // rdx = ps->r3

        "jmp  * 8*8(%%rcx)" // jump to ps->pc

        "from_native_bridge:"

        "mov  %%rdi, (%%rcx)" // ps->r0 = rdi
        "mov  %%rax, 1*8(%%rcx)" // ps->r1 = rax
        "mov  %%rbx, 2*8(%%rcx)" // ps->r2 = rbx
        "mov  %%rdx, 3*8(%%rcx)" // ps->r3 = rdx

        // recover destination control point in ps->pc
        "pop  %%rax"
        "add  $-3, %%rax" // rax = destination ctrl pt
        "mov  %%rax, 8*8(%%rcx)" // ps->pc = rax

        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  store_self_register"
        "pop  %%rsi" // handle closures
        "add  $-6, %%rsi"

        "store_self_register:"
        "mov  %%rsi, 4*8(%%rcx)" // ps->r4 = rsi
        "mov  %%rsp, 5*8(%%rcx)" // ps->fp = rsp
        "mov  %%rbp, 6*8(%%rcx)" // ps->hp = rbp

        "mov  -2*8(%%rcx), %%rsp" // restore C sp

        : // no outputs
        : // inputs
        "m" (ps)
        : // clobbers
        "rdi", "rax", "rbx", "rdx", "rsi",
        "rcx", "rbp"
    );
}
```

# Implementation of Bridges

x86-64

```
void to_native_bridge(processor_state_struct *ps) {
    __asm__ __volatile__ (

        "mov  %0, %%rcx" // copy ps into %rcx

        "mov  %%rsp, -2*8(%%rcx)" // save C sp

        // setup handler for returning from native code
        "lea  from_native_bridge(%%rip), %%rax"
        "mov  %%rax, -1*8(%%rcx)" // setup handler

        // setup frame pointer and heap pointer registers
        "mov  5*8(%%rcx), %%rsp" // rsp = ps->fp
        "mov  6*8(%%rcx), %%rbp" // rbp = ps->hp

        // setup self register
        "mov  4*8(%%rcx), %%rsi" // rsi = ps->r4

        "mov  8*8(%%rcx), %%rax" // rax = ps->pc
        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  setup_other_registers"
        "add  $3, %%rsi" // handle closures
        "push %%rsi"
        "add  $-3, %%rsi"

        "setup_other_registers:"
        "mov  (%%rcx), %%rdi" // rdi = ps->r0
        "mov  1*8(%%rcx), %%rax" // rax = ps->r1
        "mov  2*8(%%rcx), %%rbx" // rbx = ps->r2
        "mov  3*8(%%rcx), %%rdx" // rdx = ps->r3

        "jmp  * 8*8(%%rcx)" // jump to ps->pc

        "from_native_bridge:"

        "mov  %%rdi, (%%rcx)" // ps->r0 = rdi
        "mov  %%rax, 1*8(%%rcx)" // ps->r1 = rax
        "mov  %%rbx, 2*8(%%rcx)" // ps->r2 = rbx
        "mov  %%rdx, 3*8(%%rcx)" // ps->r3 = rdx

        // recover destination control point in ps->pc
        "pop  %%rax"
        "add  $-3, %%rax" // rax = destination ctrl pt
        "mov  %%rax, 8*8(%%rcx)" // ps->pc = rax

        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  store_self_register"
        "pop  %%rsi" // handle closures
        "add  $-6, %%rsi"

        "store_self_register:"
        "mov  %%rsi, 4*8(%%rcx)" // ps->r4 = rsi
        "mov  %%rsp, 5*8(%%rcx)" // ps->fp = rsp
        "mov  %%rbp, 6*8(%%rcx)" // ps->hp = rbp

        "mov  -2*8(%%rcx), %%rsp" // restore C sp

        : // no outputs
        : // inputs
        "m" (ps)
        : // clobbers
        "%rdi", "%rax", "%rbx", "%rdx", "%rsi",
        "%rcx", "%rbp"
    );
}
```

to\_native\_bridge =  
18 machine instructions

from\_native\_bridge =  
15 machine instructions

# Implementation of Bridges

x86-64

```
void to_native_bridge(processor_state_struct *ps) {
    __asm__ __volatile__ (

        "mov  %0, %%rcx" // copy ps into %rcx

        "mov  %%rsp, -2*8(%%rcx)" // save C sp

        // setup handler for returning from native code
        "lea  from_native_bridge(%%rip), %%rax"
        "mov  %%rax, -1*8(%%rcx)" // setup handler

        // setup frame pointer and heap pointer registers
        "mov  5*8(%%rcx), %%rsp" // rsp = ps->fp
        "mov  6*8(%%rcx), %%rbp" // rbp = ps->hp

        // setup self register
        "mov  4*8(%%rcx), %%rsi" // rsi = ps->r4

        "mov  8*8(%%rcx), %%rax" // rax = ps->pc
        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  setup_other_registers"
        "add  $3, %%rsi" // handle closures
        "push %%rsi"
        "add  $-3, %%rsi"

        "setup_other_registers:"
        "mov  (%%rcx), %%rdi" // rdi = ps->r0
        "mov  1*8(%%rcx), %%rax" // rax = ps->r1
        "mov  2*8(%%rcx), %%rbx" // rbx = ps->r2
        "mov  3*8(%%rcx), %%rdx" // rdx = ps->r3

        "jmp  * 8*8(%%rcx)" // jump to ps->pc

        "from_native_bridge:"

        "mov  %%rdi, (%%rcx)" // ps->r0 = rdi
        "mov  %%rax, 1*8(%%rcx)" // ps->r1 = rax
        "mov  %%rbx, 2*8(%%rcx)" // ps->r2 = rbx
        "mov  %%rdx, 3*8(%%rcx)" // ps->r3 = rdx

        // recover destination control point in ps->pc
        "pop  %%rax"
        "add  $-3, %%rax" // rax = destination ctrl pt
        "mov  %%rax, 8*8(%%rcx)" // ps->pc = rax

        "cmpq  $0x100000, -1-2*8(%%rax)" // closure?
        "jnl  store_self_register"
        "pop  %%rsi" // handle closures
        "add  $-6, %%rsi"

        "store_self_register:"
        "mov  %%rsi, 4*8(%%rcx)" // ps->r4 = rsi
        "mov  %%rsp, 5*8(%%rcx)" // ps->fp = rsp
        "mov  %%rbp, 6*8(%%rcx)" // ps->hp = rbp

        "mov  -2*8(%%rcx), %%rsp" // restore C sp

        : // no outputs
        : // inputs
        "m" (ps)
        : // clobbers
        "%rdi", "%rax", "%rbx", "%rdx", "%rsi",
        "%rcx", "%rbp"
    );
}
```

to\_native\_bridge =  
18 machine instructions

PLUS ONE  
ITERATION OF  
THE  
TRAMPOLINE

from\_native\_bridge =  
15 machine instructions

# Evaluation

# Cost of Jumps

- Compile this Scheme code using C and Native backends and vary where **dec** procedure is defined (locally or in runtime system compiled by the C backend)
- 4 possible combinations exercise cross jumps and local jumps (purely in C code, or purely in Native code)

```
(declare (standard-bindings) (not safe))

(define (dec x)      ;; variant of program defines
  (fx- x 1))        ;; dec in the runtime library

(define (run n f)
  (let loop ((n n))
    (if (fx> n 0)
        (loop (f n))))) ;; call/return dec

(time (run 100000000 dec))
```

# Cost of Jumps

cross jumps

	Native to Native		C to C (single host)		C to C (multiple hosts)		Native to C to Native	
x86-64 (i7)	0.107	<b>0.07x</b>	0.238	<b>0.17x</b>	1.436	<b>1x</b>	2.264	<b>1.58x</b>
x86-32 (P3)	2.947	<b>0.17x</b>	6.787	<b>0.40x</b>	16.854	<b>1x</b>	19.605	<b>1.16x</b>
ARM (rpi3)	2.758	<b>0.24x</b>	4.054	<b>0.35x</b>	11.431	<b>1x</b>	11.680	<b>1.02x</b>

Execution time in seconds and time relative to C backend with pure trampoline

# Cost of Jumps

cross jumps

	Native to Native		C to C (single host)		C to C (multiple hosts)		Native to C to Native
x86-64 (i7)	0.107	<b>0.07x</b>	0.238	<b>0.17x</b>	1.436	<b>1x</b>	2.264 <b>1.58x</b>
x86-32 (P3)	2.947	<b>0.17x</b>	6.787	<b>0.40x</b>	16.854	<b>1x</b>	19.605 <b>1.16x</b>
ARM (rpi3)	2.758	<b>0.24x</b>	4.054	<b>0.35x</b>	11.431	<b>1x</b>	11.680 <b>1.02x</b>

Execution time in seconds and time relative to C backend with pure trampoline

- Cross jumps cost **2%** to **58%** more than pure trampoline



# Cost of Jumps

cross jumps

	Native to Native		C to C (single host)		C to C (multiple hosts)		Native to C to Native	
x86-64 (i7)	0.107	<b>0.07x</b>	0.238	<b>0.17x</b>	1.436	<b>1x</b>	2.264	<b>1.58x</b>
x86-32 (P3)	2.947	<b>0.17x</b>	6.787	<b>0.40x</b>	16.854	<b>1x</b>	19.605	<b>1.16x</b>
ARM (rpi3)	2.758	<b>0.24x</b>	4.054	<b>0.35x</b>	11.431	<b>1x</b>	11.680	<b>1.02x</b>

Execution time in seconds and time relative to C backend with pure trampoline

- Local jumps within code generated by Native backend are up to **14x** faster than inter host trampoline and **2.4x** faster than intra host trampoline

# Conclusion

# Desiderata

1. **Cross jumps are transparent** (All control points support both the trampoline and *branch-to-address*)
2. **Native code jumps are plain *branch-to-address* instructions** (The single machine jump instruction is optimal)
3. **Cost of cross jumps is similar to C backend local jumps** (Only 1.58x slower for x86-64, 1.16x for x86-32 and 1.02x for ARM)

# Use cases

- 1. Useful tool for native backend development  
(all system features available throughout  
development)**
- 2. Possibly completely avoid reimplementing the  
I/O and other complex subsystems**



# Motivation for Native Backend

chez-9.5.1-m64	1111111111111111111111111111222222222223333333333444444446778
stalin-unknown	1111111111111111222222333445566667
gambit/gerbil	1111122222222222222222223333333333333333444444555556789
mit-9.2.1	11123344445566666666667777788888899
ypsilon-unknown	114568999
bigloo-4.3a	122222222233333344445555556666777777777888899
racket-7.0/r7rs	12233333334444444444444555555555566666666677777889
cyclone-0.9.2	122335566777777888888899999999
bones-unknown	123445566677888888999999999
femtolisp-unknown	135569
picrin-unknown	19
rhizome-unknown	19
larceny-1.3	22222333333334444444444445555555555555666667778999
petite-9.5.1-m64	234455567777888899999
guile-2.2.4	256677777888889999999999
chicken-4.13.0	445555566666666777777888888889999
kawa-3.0	46679
rscheme-unknown	4
gauche-0.9.6	66677788889999
chickensci-4.13.0	7
sagittarius-0.9.2	888899
chibi-unknown	8

# Motivation for Native Backend

mostly  
loops

	C Backend	x86-64 Backend	Chez Scheme 9.5.1 --optimize-level 3
ack	12.21 <b>1.11x</b>	10.96 <b>1x</b>	13.73 <b>1.25x</b>
fib	3.25 <b>1.26x</b>	2.57 <b>1x</b>	5.28 <b>2.05x</b>
tak	8.06 <b>1.10x</b>	7.31 <b>1x</b>	9.56 <b>1.31x</b>
takl	3.40 <b>0.90x</b>	3.78 <b>1x</b>	8.15 <b>2.16x</b>
sum	2.85 <b>0.56x</b>	5.12 <b>1x</b>	11.76 <b>2.30x</b>

Unsafe fixnum benchmarks with frequent jumps

(Preliminary results)